

# Homework 4:

# Reinforcement Learning

# Report

110652019 林楷傑

## Part I. Implementation (-5 if not explain in detail):

- **Part1**

```
# Begin your code
# TODO
"""
Use epsilon greedy select action
"""
if random.random() > self.epsilon:
    action = np.argmax(self.qtable[state])
else:
    action = random.randint(0, self.env.action_space.n-1)
return action
# End your code
```

```
# Begin your code
# TODO
"""
update Q table by the formula:  $Q(s,a) = (1-\alpha) * Q(s,a) + \alpha * (r + \gamma * \max(Q(s')))$ 
"""
self.qtable[state][action] = (1-self.learning_rate) * self.qtable[state][action] + self.learning_rate*(reward + self.gamma * max(self.qtable[next_state]))
# End your code
```

```
# Begin your code
# TODO
"""
Return the max Q value of the giving state
"""
return max(self.qtable[state])
# End your code
```

- **Part2**

```
# Begin your code
# TODO
"""
Use numpy function linspace to slice the interval into num_bins part
"""
bins = np.linspace(lower_bound, upper_bound, num_bins)
return bins[1:-1]
# End your code
```

```
# Begin your code
# TODO
"""
Use numpy function digitize to discretize the value with given bins.
"""
bin_index = np.digitize(value, bins)
return bin_index-1
# End your code
```

```
# Begin your code
# TODO
"""
discretize the giving state
"""
return [self.discretize_value(observation[i], self.bins[i]) for i in range(4)]
# End your code
```

```
# Begin your code
# TODO
"""
Use epsilon greedy select action
"""
if random.random() > self.epsilon:
    return np.argmax(self.qtable[state[0]][state[1]][state[2]][state[3]])
return np.array(random.randrange(self.n_actions))
# End your code
```

```
# Begin your code
# TODO
"""
update Q table by the formula:  $Q(s,a) = (1-\alpha) * Q(s,a) + \alpha * (r + \gamma * \max_{a'} Q(s',a'))$ 
"""
if done:
    self.qtable[state[0]][state[1]][state[2]][state[3]][action] = self.qtable[state[0]][state[1]][state[2]][state[3]][action] + self.learning_rate * (
        reward - self.qtable[state[0]][state[1]][state[2]][state[3]][action])
else:
    self.qtable[state[0]][state[1]][state[2]][state[3]][action] = self.qtable[state[0]][state[1]][state[2]][state[3]][action] + self.learning_rate * (
        reward + self.gamma * np.max(self.qtable[next_state[0]][next_state[1]][next_state[2]][next_state[3]]) - self.qtable[state[0]][state[1]][state[2]][state[3]][action])
# End your code
```

```
# Begin your code
# TODO
"""
Return max Q value of giving state
"""
state = self.discretize_observation(self.env.reset())
return np.max(self.qtable[state[0]][state[1]][state[2]][state[3]])
# End your code
```

## Part3

```
# Begin your code
# TODO
"""
s: state; s':next state; a: action; r: reward
Sample a batch data from memory buffer.
Q_evaluate: use torch gather to get Q(s,a)
Q_next: use target net comput Q'(s', )
Q_target:  $r + \gamma * Q'(s', \arg\max_{a'} Q'(s', a))$  * (1-done)
Use mean square loss between Q and Q_target to update the evaluate net
"""
batch = self.buffer.sample(self.batch_size)
#batch = (observations, actions, rewards, next_observations, done)
state_batch = Variable(Tensor(np.array(batch[0])))
action_batch = Variable(Tensor(batch[1]))
reward_batch = Variable(Tensor(batch[2]))
mask_batch = Variable(Tensor(batch[4]))
next_state_batch = Variable(Tensor(np.array(batch[3])))

Q_evaluate = torch.gather(self.evaluate_net(state_batch), 1, action_batch.view(-1, 1).long())
Q_next = self.target_net(next_state_batch).detach()
Q_target = reward_batch.view(-1, 1) + self.gamma * Q_next.max(1)[0].view(self.batch_size, 1) * (1-mask_batch.view(self.batch_size, 1))
loss = F.mse_loss(Q_evaluate, Q_target)

self.optimizer.zero_grad()
loss.backward()
self.optimizer.step()

# End your code
```

```

# Begin your code
# TODO
"""
Use epsilon greedy select action
"""
state = torch.Tensor(state)
if random.uniform(0,1) > self.epsilon:
    Q_val = self.evaluate_net(state)
    action = torch.argmax(Q_val).item()
else:
    action = np.array(random.randrange(self.n_actions))
# End your code

```

```


# Begin your code
# TODO
"""
Return the max Q that initial state pass through the target net.
"""
return max(self.target_net(Tensor(self.env.reset()))))
# End your code

```


## Part II. Experiment Results:

Please paste [taxi.png](#), [cartpole.png](#), [DQN.png](#) and [compare.png](#) here.


### 1. taxi.png:




## 2. cartpole.png



## 3. DQN.png



#### 4. compare.png



### Part III. Question Answering (50%):

1. Calculate the optimal Q-value of a given state in Taxi-v3, and compare with the Q-value you learned (Please screenshot the result of the “`check_max_Q`” function to show the Q-value you learned). (10%)

```
# Q opt: for state 243, at least need 9 step to reach the goal.
print("Q opt:", -(1-self.gamma**9)/(1-self.gamma) + 20 * self.gamma**9)
return max(self.qtable[state])
# End your code
```

```
average reward: 7.56
Initial state:
taxi at (2, 2), passenger at Y, destination at R
Q opt: 1.6226146700000017
max Q: 1.6226146699999995
```

2. Calculate the max Q-value of the initial state in CartPole-v0, and compare with the Q-value you learned. (Please screenshot the result of the “`check_max_Q`” function to show the Q-value you learned) (10%)

```
# Q opt: Need at least 200 step to reach the final goal. The optimal reward is 200.
print("Q opt: ", (1-self.gamma**200)/(1-self.gamma))
state = self.discretize_observation(self.env.reset())
return np.max(self.qtable[state[0]][state[1]][state[2]][state[3]])
# End your code
```

```
average reward: 196.82
Q opt: 33.25795863300011
max Q: 30.639871238426274
```

3.
  - a. Why do we need to discretize the observation in Part 2? (3%)  
Since the observation of cartpole environment is continuous. We need to discretize the observation to present on the Q table.
  - b. How do you expect the performance will be if we increase “`num_bins`”? (3%)

The run time of the program will be longer, but the performance will be better. Since the agent can gain more knowledge from the environment.

- c. Is there any concern if we increase “num\_bins”? (3%)

The runtime of the program would be longer.

4. Which model (DQN, discretized Q learning) performs better in Cartpole-v0, and what are the reasons? (5%)

DQN. The neural network can deal the continuous space. It solves the problem that discretized Q learning encountered, that is continuous space is hard to presented on the tabular method. Thus, DQN would perform better.

5.

- a. What is the purpose of using the epsilon greedy algorithm while choosing an action? (3%)

For exploration, use epsilon greedy algorithm to make the agent fully explore the environment.

- b. What will happen, if we don't use the epsilon greedy algorithm in the CartPole-v0 environment? (3%)

The policy will stuck in a local optimal action since the agent will always choose the max Q which first update previously and is not actually optimal. The agent is lack of exploration.

- c. Is it possible to achieve the same performance without the epsilon greedy algorithm in the CartPole-v0 environment? Why or Why not? (3%)

No, without exploration, the agent can not fully explore actions that is likely to be better.

- d. Why don't we need the epsilon greedy algorithm during the testing section? (3%)

Since we have successfully trained out agent, the agent can perform well policy. We do not need randomness to disturb its decision.

6. Why does “`with torch.no_grad():`” do inside the “choose\_action” function in DQN? (4%)

`with torch.no_grad():` means the back propagation will not propagate through the computation. Since the decision making is depend on choosing the maximum Q value, we don't need to tack the gradient of this process.